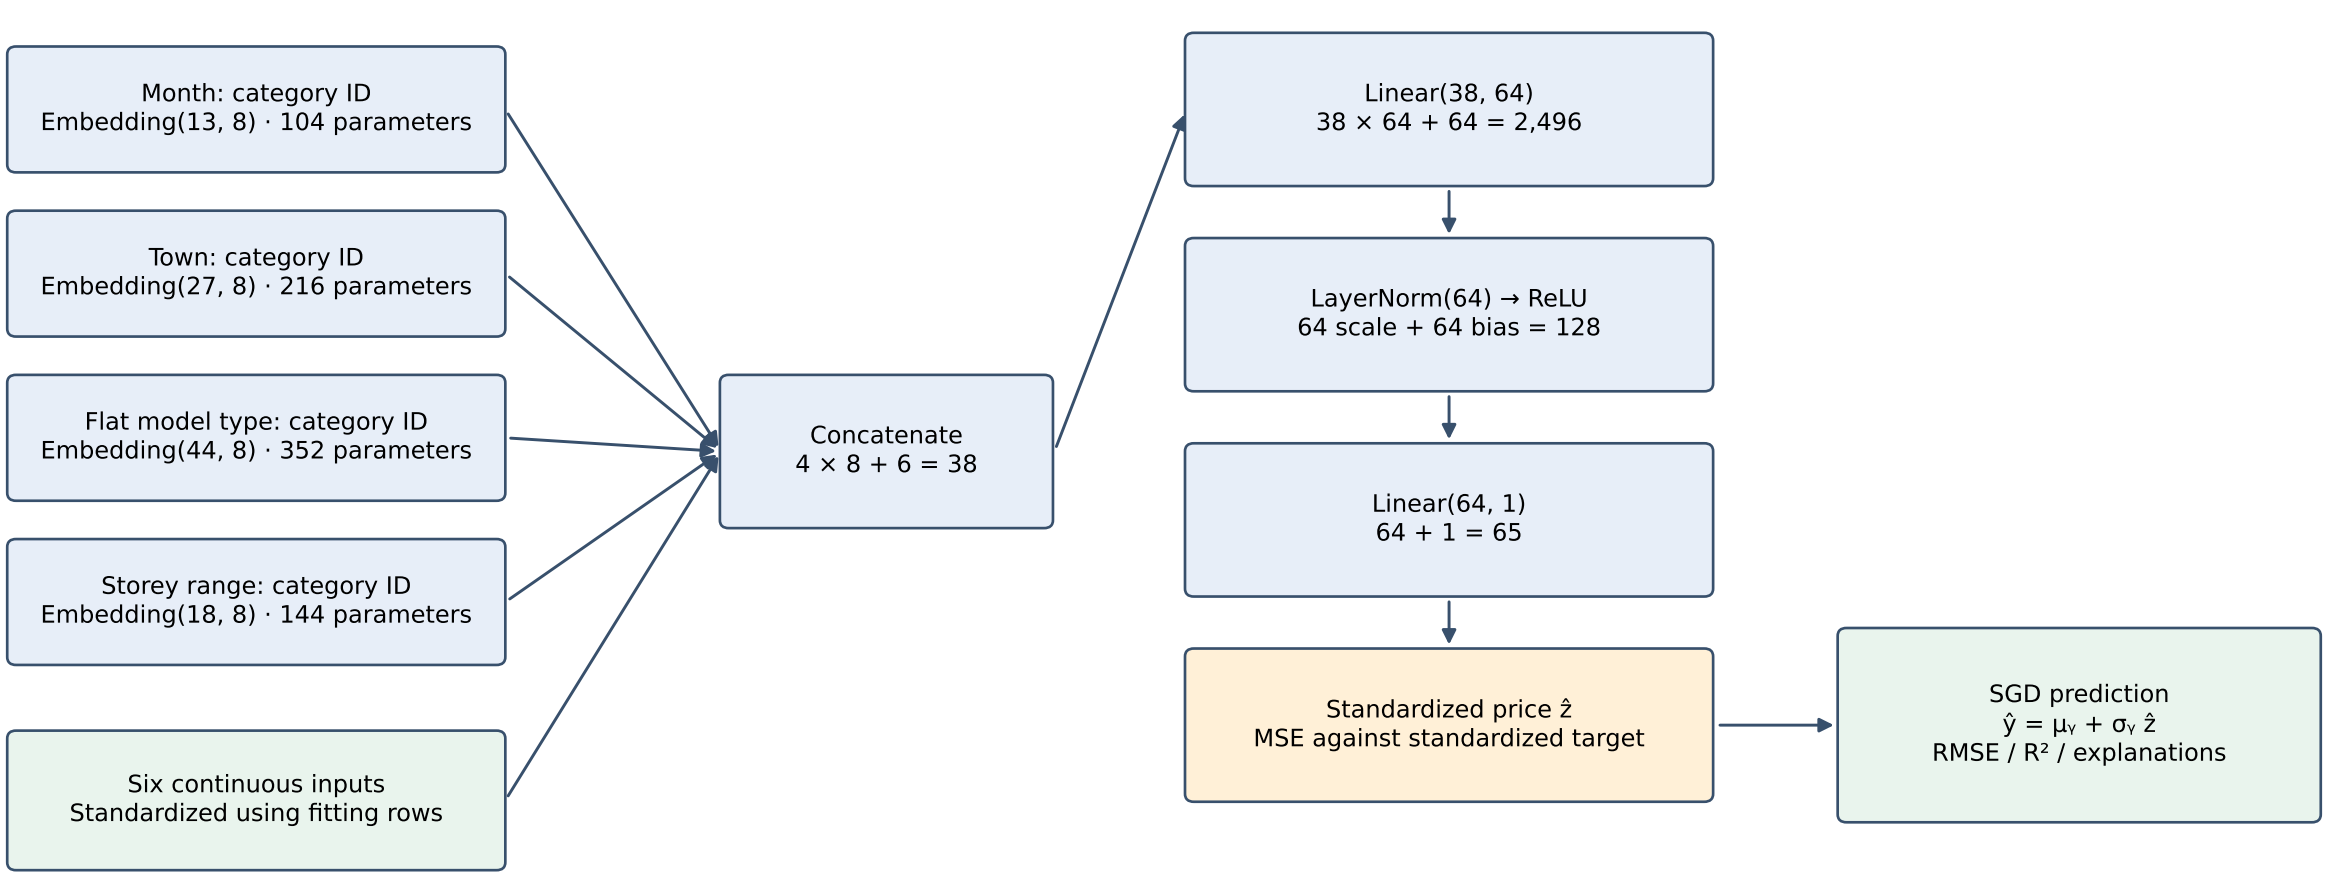
Supplied baseline: four independent embeddings + LayerNorm MLP
Month: category ID
Embedding(13, 8) · 104 parameters
Town: category ID
Embedding(27, 8) · 216 parameters
Flat model type: category ID
Embedding(44, 8) · 352 parameters
Storey range: category ID
Embedding(18, 8) · 144 parameters
Six continuous inputs
Standardized using fitting rows
Concatenate
4 × 8 + 6 = 38
Linear(38, 64)
38 × 64 + 64 = 2,496
LayerNorm(64) → ReLU
64 scale + 64 bias = 128
Linear(64, 1)
64 + 1 = 65
Standardized price ẑ
MSE against standardized target
SGD prediction
ŷ = μY + σY ẑ
RMSE / R² / explanations
Trainable parameters: 816 embeddings + 2,496 hidden linear + 128 LayerNorm + 65 output = 3,505